

Bodyfat Prediction

Yilun Chen
Jiyun Chen
Huiyu Jiang

University of Wisconsin-Madison

February 7, 2018


Outline

- Background and Data Description
- Model Selection and Data Processing
- Variable Selection
- Model Diagnostic and Prediction


Background

- Motivation:
Use data to train models to make predictions.
- Target:
A clean and robust model


Model Selection




(a) age



(b) weight



(c) height



(d) abdomen

Data Processing

- false record:

| | BODYFAT | AGE | WEIGHT | HEIGHT | ADIPOSIITY |
|-----|---------|-----|--------|--------|------------|
| 182 | 0 | 40 | 118.5 | 68 | 18.1 |


- check outliers:

| | BODYFAT | AGE | WEIGHT | HEIGHT | ADIPOSIITY |
|----|---------|-----|--------|--------|------------|
| 42 | 31.7 | 44 | 205 | 29.5 | 29.9 |


(e) 42

| | BODYFAT | AGE | WEIGHT | HEIGHT | ADIPOSIITY |
|----|---------|-----|--------|--------|------------|
| 39 | 33.8 | 46 | 363.15 | 72.25 | 48.9 |

(f) 39



(g)



(h)

Variable Selection

- AIC Stepwise:

```
Call:
lm(formula = BODYFAT ~ ABDOMEN + WEIGHT + WRIST + BICEPS + AGE,
    data = data_body)


Residuals:
    Min       1Q   Median       3Q      Max
-9.4993 -2.6962 -0.2706  3.1972  8.7266

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -23.71669     7.06737  -3.356 0.000918 ***
ABDOMEN       0.83796     0.06314  13.272 < 2e-16 ***
WEIGHT      -0.08539     0.03055  -2.795 0.005607 **
WRIST       -1.65891     0.45052  -3.682 0.000285 ***
BICEPS       0.27709     0.13981   1.982 0.048619 *
AGE          0.03881     0.02615   1.484 0.139047

---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 3.95 on 243 degrees of freedom
Multiple R-squared:  0.7356,    Adjusted R-squared:  0.7301
F-statistic: 135.2 on 5 and 243 DF,  p-value: < 2.2e-16
```

Variable Selection



Variable Selection

- elastic net:



Variable Selection

| | |
|----------------|---------------|
| WEIGHT | 4.1483 |
| ABDOMEN | 4.1483 |

Call:

```
lm(formula = BODYFAT ~ WEIGHT + ABDOMEN, data = df)
```

Residuals:

| | Min | 1Q | Median | 3Q | Max |
|--|----------|---------|---------|--------|--------|
| | -10.0573 | -2.9614 | -0.1035 | 2.9452 | 9.9097 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|-----------|------------|---------|--------------|
| (Intercept) | -42.28955 | 2.45945 | -17.20 | < 2e-16 *** |
| WEIGHT | -0.12314 | 0.01955 | -6.30 | 1.36e-09 *** |
| ABDOMEN | 0.90011 | 0.05179 | 17.38 | < 2e-16 *** |


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.054 on 246 degrees of freedom


Multiple R-squared: 0.718, Adjusted R-squared: 0.7157

F-statistic: 313.2 on 2 and 246 DF, p-value: < 2.2e-16


Diagnostic




(i) resid



(j) QQ-plot



(k) std.res



(l) cook distance

Prediction

